



Learn Git and GitHub without any code!

Using the Hello World guide, you'll start a branch, write comments, and open a pull request.

Read the guide

JosephDenney / KingCountyRealEstate

<> Code

! Issues

🔗 Pull requests

▶ Actions

📁 Projects

📖 Wiki

🛡 Security

🔗 master ▾



KingCountyRealEstate / Mod2ProjectEDA.ipynb



JosephDenney sunday final nb

🕒 History

👤 1 contributor



Raw

Blame



987 lines (987 sloc) 437 KB

Housing Analysis in King County, Washington

EDA, data cleaning, feature engineering notebook

```
In [24]: import warnings

warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
import csv

import scipy.stats as scs
import statsmodels.api as sm
import statsmodels.formula.api as sms
import scipy.stats as stats
from statsFunctions import check_model as sf
from pltfunctions import hist_kde_plots
from haversine import haversine
from math import sqrt

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import f_regression
import sklearn.metrics as metrics

import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv(r'data\kc_house_data.csv')
```

Data description as follows

Variable	Description
Id	Unique ID for each home sold
Date	Date of the home sale
Price	Price of each home sold
Bedrooms	Number of bedrooms
Bathrooms	Number of bathrooms, where .5 accounts for a room with a toilet but no shower
Sqft_living	Square footage of the apartments interior living space
Sqft_lot	Square footage of the land space
Floors	Number of floors
Waterfront	A dummy variable for whether the apartment was overlooking the waterfront or not
View	An index from 0 to 4 of how good the view of the property was

Condition	An index from 1 to 5 on the condition of the apartment,
Grade	An index from 1 to 13, where 1-3 falls short of building construction and design, 7 has an average level of construction and design, and 11-13 have a high quality level of construction and design
Sqft_above	The square footage of the interior housing space that is above ground level
Sqft_basement	The square footage of the interior housing space that is below ground level
Yr_built	The year the house was initially built
Yr_renovated	The year of the house's last renovation
Zipcode	What zipcode area the house is in
Lat	Lattitude
Long	Longitude
Sqft_living15	The square footage of interior housing living space for the nearest 15 neighbors
Sqft_lot15	The square footage of the land lots of the nearest 15 neighbors

In [25]: *df.info() # a good initial picture of the data*
initial thoughts for necessary data and a regression upon seeing the data -
price is y target
1) Unique identifiers (column= id) are unnecessary for a regression
2) lat and long wont be needed if include zipcode and will just be noise (decided to keep lat/long over zip)
3) square footage and home quality is likely to change based on location
4) anticipate autocorrelation between location and home features - dropping zipcode and keeping lat long
5) date - as time goes on, home prices are likely to go up - if the sales are all within a close time frame then we can ignore date
6) nearest 15 neighbors data - will autocorrelate with homes nearby, remove data
7) already have square footage of home, remove above and below ground sqft

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               21597 non-null  int64
1   date            21597 non-null  object
2   price           21597 non-null  float64
3   bedrooms        21597 non-null  int64
4   bathrooms       21597 non-null  float64
5   sqft_living     21597 non-null  int64
6   sqft_lot        21597 non-null  int64
7   floors          21597 non-null  float64
8   waterfront      19221 non-null  float64
9   view            21534 non-null  float64
10  condition       21597 non-null  int64
11  grade           21597 non-null  int64
12  sqft_above      21597 non-null  int64
```

```

13 sqft_basement    21597 non-null    object
14 yr_built         21597 non-null    int64
15 yr_renovated     17755 non-null    float64
16 zipcode          21597 non-null    int64
17 lat              21597 non-null    float64
18 long             21597 non-null    float64
19 sqft_living15    21597 non-null    int64
20 sqft_lot15       21597 non-null    int64

```

dtypes: float64(8), int64(11), object(2)

memory usage: 3.5+ MB

In [26]: *# before we drop any null values (in yr_renovated and waterfront specifically) we should address the issues above.*

```
df_init = df.drop(columns=['id','sqft_living15','sqft_above','sqft_base
ment'])
```

```
# filter for more recent home sales
```

```
df_init['date'].unique()
```

```
# can remove date and null values
```

```
print(df['yr_renovated'].unique())
```

```
print(df.groupby(df['yr_renovated']).count())
```

```
# given that these dates are very spread out and also that 17011 of the
homes sold out have no renovation year at all, I will remove the data f
or the purposes of a linear regression
```

```

[ 0. 1991.   nan 2002. 2010. 1992. 2013. 1994. 1978. 2005. 2003. 198
4.
1954. 2014. 2011. 1983. 1945. 1990. 1988. 1977. 1981. 1995. 2000. 199
9.
1998. 1970. 1989. 2004. 1986. 2007. 1987. 2006. 1985. 2001. 1980. 197
1.
1979. 1997. 1950. 1969. 1948. 2009. 2015. 1974. 2008. 1968. 2012. 196
3.
1951. 1962. 1953. 1993. 1996. 1955. 1982. 1956. 1940. 1976. 1946. 197
5.
1964. 1973. 1957. 1959. 1960. 1967. 1965. 1934. 1972. 1944. 1958.]

```

```

          id  date  price  bedrooms  bathrooms  sqft_living  sq
ft_lot \
yr_renovated
0.0          17011  17011  17011      17011      17011      17011
17011
1934.0          1      1      1          1          1          1
1
1940.0          2      2      2          2          2          2
2
1944.0          1      1      1          1          1          1
1
1945.0          3      3      3          3          3          3
3
...          ...      ...      ...      ...      ...      ...
...
2011.0          9      9      9          9          9          9
9
2012.0          8      8      8          8          8          8
8
2013.0         31     31     31         31         31         31
31
2014.0         73     73     73         73         73         73
73

```

```

/>
2015.0      14      14      14      14      14      14
14

      floors  waterfront  view  condition  grade  sqft_above
\
yr_renovated
0.0      17011      15157  16961      17011  17011      17011
1934.0      1      1      1      1      1      1
1940.0      2      2      2      2      2      2
1944.0      1      1      1      1      1      1
1945.0      3      2      3      3      3      3
...      ...      ...      ...      ...      ...      ...
2011.0      9      7      9      9      9      9
2012.0      8      7      8      8      8      8
2013.0     31     29     31     31     31     31
2014.0     73     64     73     73     73     73
2015.0     14     13     14     14     14     14

      sqft_basement  yr_built  zipcode  lat  long  sqft_livi
ng15 \
yr_renovated
0.0      17011      17011      17011  17011  17011      1
7011
1934.0      1      1      1      1      1
1
1940.0      2      2      2      2      2
2
1944.0      1      1      1      1      1
1
1945.0      3      3      3      3      3
3
...      ...      ...      ...      ...      ...
...
2011.0      9      9      9      9      9
9
2012.0      8      8      8      8      8
8
2013.0     31     31     31     31     31
31
2014.0     73     73     73     73     73
73
2015.0     14     14     14     14     14
14

      sqft_lot15
yr_renovated
0.0      17011
1934.0      1
1940.0      2
1944.0      1
1945.0      3
...      ...
2011.0      9
2012.0      8
2013.0     31
2014.0     73
2015.0     14

```

[70 rows x 20 columns]

```
In [27]: df_init = df_init.dropna() # drop null values
df_init = df_init.drop(columns=['date', 'yr_renovated']) # drop date column
```

```
In [28]: df_init.head()
```

```
Out[28]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
1	538000.0	3	2.25	2570	7242	2.0	0.0	0.0
3	604000.0	4	3.00	1960	5000	1.0	0.0	0.0
4	510000.0	3	2.00	1680	8080	1.0	0.0	0.0
5	1230000.0	4	4.50	5420	101930	1.0	0.0	0.0
6	257500.0	3	2.25	1715	6819	2.0	0.0	0.0

```
In [29]: df_init.info() # data is all in integer or float dtype
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 15762 entries, 1 to 21596
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   price           15762 non-null  float64
1   bedrooms        15762 non-null  int64
2   bathrooms       15762 non-null  float64
3   sqft_living     15762 non-null  int64
4   sqft_lot        15762 non-null  int64
5   floors          15762 non-null  float64
6   waterfront      15762 non-null  float64
7   view            15762 non-null  float64
8   condition       15762 non-null  int64
9   grade           15762 non-null  int64
10  yr_built        15762 non-null  int64
11  zipcode         15762 non-null  int64
12  lat             15762 non-null  float64
13  long            15762 non-null  float64
14  sqft_lot15      15762 non-null  int64
dtypes: float64(7), int64(8)
memory usage: 1.9 MB
```



```
In [30]: features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'sqft_lot15', 'floors', 'waterfront', 'view', 'condition', 'grade', 'yr_built', 'yr_renovated', 'lat', 'long']
```

```
In [9]: for feature in features:
        hist_kde_plots(feature, 'price', df_init)
```


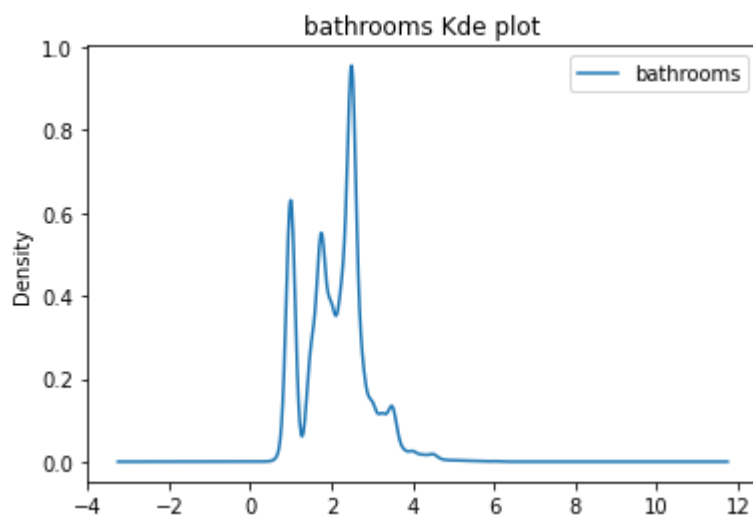
bedrooms

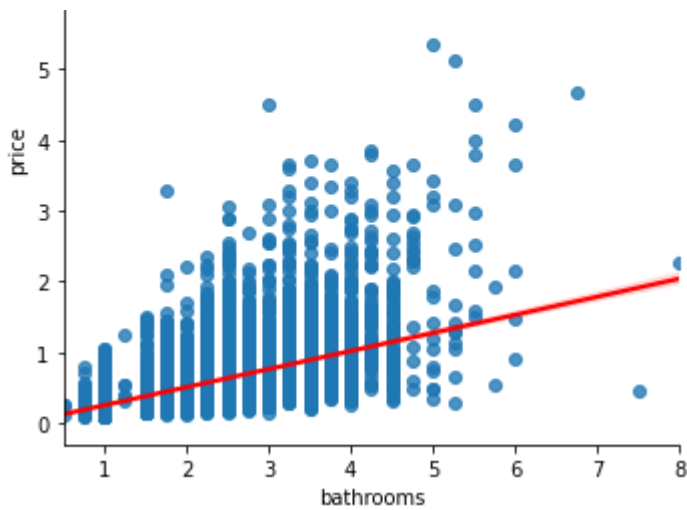
bedrooms Kde plot




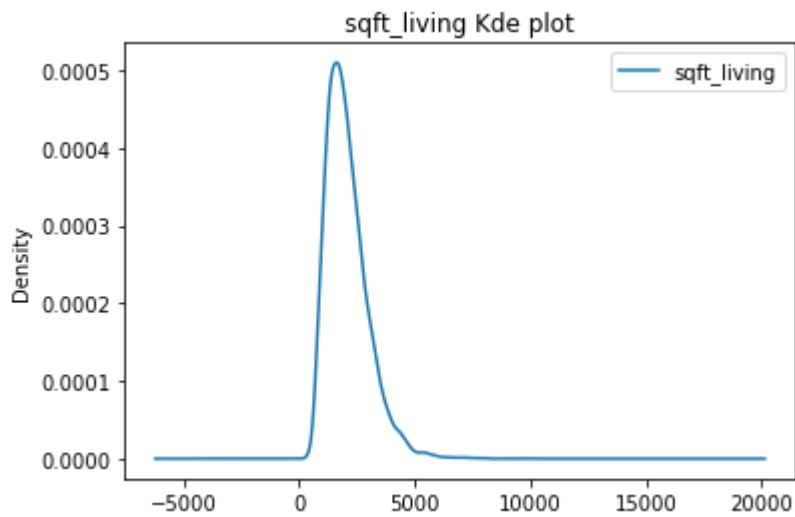


bathrooms








sqft_living

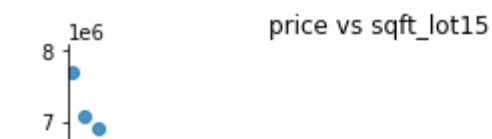




sqft_lot



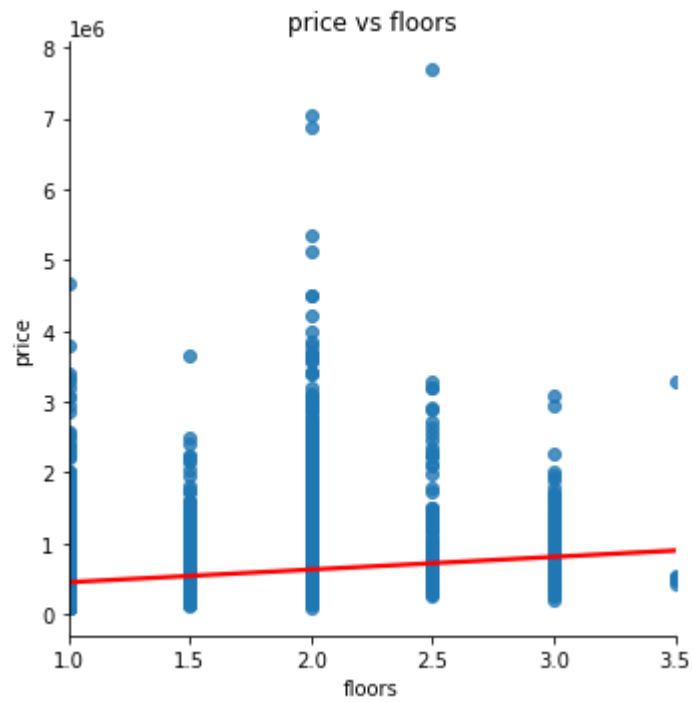
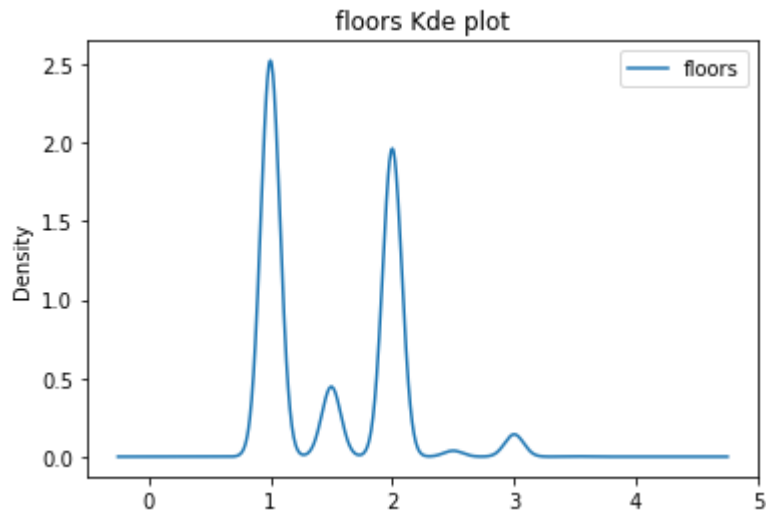


sqft_lot15



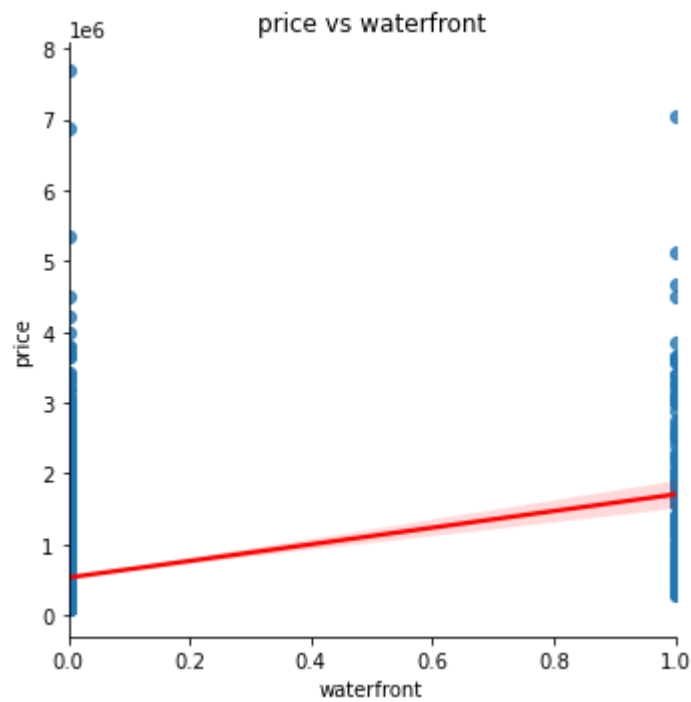
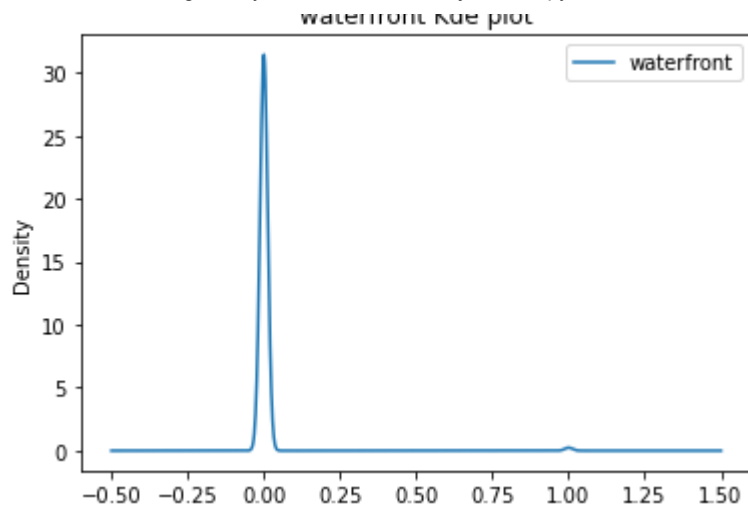


floors

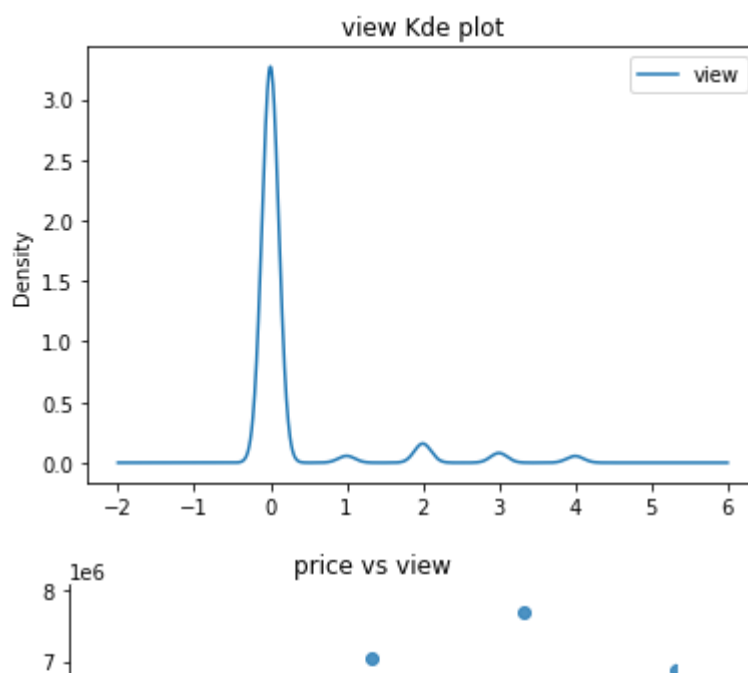



waterfront

waterfront Kde plot





view


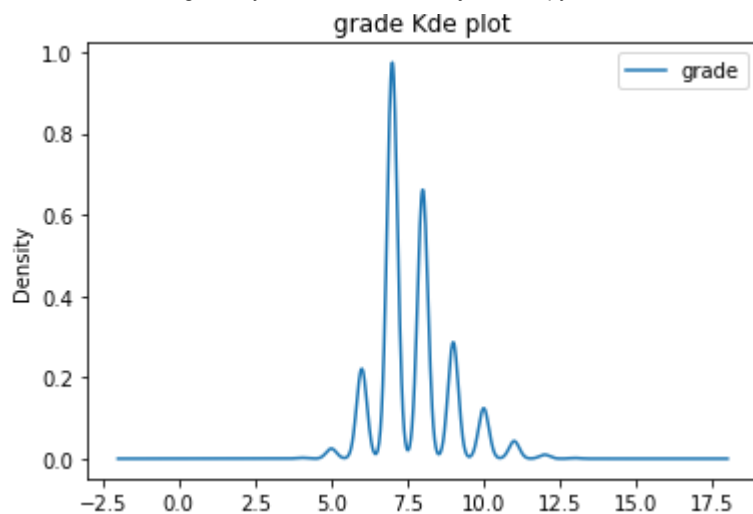





condition




grade



yr_built





yr_renovated

```
-----
----
KeyError                                Traceback (most recent call last)
~\anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method, tolerance)
    2645         try:
-> 2646             return self._engine.get_loc(key)
    2647         except KeyError:

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

KeyError: 'yr_renovated'
```

During handling of the above exception, another exception occurred:

```
KeyError                                Traceback (most recent call last)
<ipython-input-9-ca5a9d20d9bc> in <module>
      1 for feature in features:
----> 2     hist_kde_plots(feature, 'price', df_init)

~\desktop\coursework\phase_1\Phase2\Phase_2_Project\pltfunctions.py in hist_kde_plots(feature, target, df)
     52     print(feature)
     53     #kde plot
--> 54     df[feature].plot.kde(label=feature)
     55     plt.title("{} Kde plot".format(feature))
     56     plt.legend()
```

```
~\anaconda3\lib\site-packages\pandas\core\frame.py in getitem (self,
```

```

key)
2798         if self.columns.nlevels > 1:
2799             return self._getitem_multilevel(key)
-> 2800         indexer = self.columns.get_loc(key)
2801         if is_integer(indexer):
2802             indexer = [indexer]

~\anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method, tolerance)
2646         return self._engine.get_loc(key)
2647     except KeyError:
-> 2648         return self._engine.get_loc(self._maybe_cast_indexer(key))
2649     indexer = self.get_indexer([key], method=method, tolerance=tolerance)
2650     if indexer.ndim > 1 or indexer.size > 1:

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

KeyError: 'yr_renovated'

```

In [31]: *# funky data in the following columns - 1) yr_renovated - a lot of homes that haven't been renovated and they are all listed as 0's. Need to reconcile this data - because of groupby count analysis above decided to just remove this data (<17000 homes were never renovated per original dataset)*
2) remove any price data that is greater than 3 standard deviations from the mean
df_init.describe()

Out[31]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	fl
count	1.576200e+04	15762.000000	15762.000000	15762.000000	1.576200e+04	15762.000000
mean	5.413172e+05	3.378949	2.120797	2084.512372	1.528082e+04	1.528082e+04
std	3.722258e+05	0.935301	0.766772	918.617686	4.182288e+04	4.182288e+04
min	8.200000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000
25%	3.210000e+05	3.000000	1.750000	1430.000000	5.048500e+03	1.000000
50%	4.500000e+05	3.000000	2.250000	1920.000000	7.602000e+03	1.000000
75%	6.448750e+05	4.000000	2.500000	2550.000000	1.072000e+04	2.000000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.000000

In [32]: df_init.isna().sum()

```
Out[32]: price      0
          bedrooms   0
          bathrooms  0
          sqft_living 0
          sqft_lot    0
          floors      0
          waterfront  0
          view        0
          condition   0
          grade       0
          yr_built    0
          zipcode     0
          lat         0
          long        0
          sqft_lot15  0
          dtype: int64
```

```
In [33]: # remove outliers and data will be cleaned, can move on to feature engi
          # neering
          # we can standartize the data as well to do this
          # we can see that there are no homes below 3 standard deviations with t
          # he describe table
          # however there are homes above 3 standard deviations of the mean price
          z = abs(stats.zscore(df_init))
          print(z)
          threshold = 3
          print(np.where(z>3))
          avg_price = df_init['price'].mean()
          stdev_price = df_init['price'].std()
          upper_bound = avg_price + 3*stdev_price

          # df_clean = df_init[(z < 3).all(axis=1)] # this doesnt work because we
          # remove some data (waterfront) that we don't want to remove - remove onl
          # y homes sold with a price above 3 standard deviations
          df_clean = df_init[df_init['price']<upper_bound]

          [[0.00891201 0.40517583 0.16850812 ... 1.16734842 0.74967289 0.1880666
            ]
           [0.16840532 0.66403252 1.14666608 ... 0.2768393  1.27560722 0.2823963
            ]
           [0.08413755 0.40517583 0.15754454 ... 0.41567829 1.19770555 0.1929278
            5]
           ...
           [0.37966644 0.66403252 0.49456077 ... 0.34969792 1.05528338 0.2037584
            3]
           [0.37402184 1.47438418 1.78780781 ... 0.25409085 0.60752848 0.3893795
            5]
           [0.58116341 1.47438418 1.78780781 ... 0.25192673 0.60752848 0.4126134
            6]]
          (array([ 3, 3, 3, ..., 15740, 15740, 15751], dtype=int64), a
           rray([ 2, 3, 14, ..., 6, 7, 3], dtype=int64))
```

```
In [34]: df_clean.isna().sum()
```

```
Out[34]: price      0
          bedrooms   0
          bathrooms  0
          sqft_living 0
```



```

sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
grade         0
yr_built      0
zipcode       0
lat           0
long          0
sqft_lot15    0
dtype: int64

```

```

In [35]: df_clean.price.max() # have removed absurdly high priced homes from df
df_clean.waterfront.max() # where we discover we removed data that is binary with low standard deviation compared to the max, which is 1.
# Lets try a different method
df_clean.describe()

```

```

Out[35]:

```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	fl
count	1.548000e+04	15480.000000	15480.000000	15480.000000	1.548000e+04	15480
mean	5.084310e+05	3.362532	2.093169	2037.260401	1.507320e+04	1.507320
std	2.615181e+05	0.926927	0.733848	834.955963	4.123850e+04	0.926927
min	8.200000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000
25%	3.200000e+05	3.000000	1.500000	1420.000000	5.012750e+03	1.500000
50%	4.490000e+05	3.000000	2.250000	1900.000000	7.560000e+03	1.500000
75%	6.275000e+05	4.000000	2.500000	2510.000000	1.050000e+04	2.000000
max	1.650000e+06	33.000000	7.500000	7350.000000	1.651359e+06	33.000000

```

In [36]: # after having completed a regression, discovered that lat and long are
not worthwhile in a regression analysis
# as a result, will create a distance from seattle feature to create a
regression

```

```

In [85]: df_clean['lat'] = [round(i,4) for i in df_clean['lat']]
df_clean['long'] = [round(i,4) for i in df_clean['long']]

```

```

In [86]: df_clean['geo_loc'] = list(zip(df_clean['lat'], df_clean['long']))

```

```

In [87]: Seattle = [47.6219, -122.3517] # defining Seattle's location

```

```

In [88]: distance = []
for i in df_clean['geo_loc']:
    distance.append((haversine((Seattle),(i), unit='mi')))
rounded = [round(i,4) for i in distance]

df_clean['distance'] = rounded

```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 15480 entries, 1 to 21596
Data columns (total 17 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   price                 15480 non-null  float64
 1   bedrooms              15480 non-null  int64
 2   bathrooms             15480 non-null  float64
 3   sqft_living           15480 non-null  int64
 4   sqft_lot              15480 non-null  int64
 5   floors                15480 non-null  float64
 6   waterfront            15480 non-null  float64
 7   view                  15480 non-null  float64
```